

VCFloat 2.1 Reference Manual

Andrew W. Appel

Ariel E. Kellison

December 4, 2023

Contents

1	Introduction	2
2	Floating-point functional models	3
3	NaNs	3
4	Notation Scopes	4
5	Operators	4
6	Example	4
7	Reification	5
8	Boundsmap	5
9	Valmap and reflection	6
10	Real-valued functional model	6
11	Round-off theorem	7
12	Verification conditions	8
13	Letting VCFloat calculate the accuracy bound	8
14	<code>field_simplify</code>	9
15	<code>prune_terms</code>	10
16	<code>error_rewrites</code>	10
17	User-defined functions	12
18	Abstract <i>versus</i> transparent valmap	12
19	Annotations	13
20	Verified Software Toolchain	13
21	Examples	14
22	Nonstandard floating-point-like types	15
23	Functional Modeling Languages	16
24	<i>General</i> Modeling Language	17
25	<i>Standard</i> Modeling Language	18
26	Generic expressions	18
27	Equivalence relations and rewriting morphisms	19
28	Conversions between models	20

1 Introduction

VCFloat is a tool for Coq proofs about floating-point round-off error. When performing a computation such as $x \times 5.7 + y$ in floating-point with a fixed number of mantissa bits, the result of $x \times 5.7$ cannot always be represented exactly in the same number of bits, ditto the result of the addition $+y$, so some low-order bits must be thrown away—there is *round-off error*.

We can state this more formally with a bit of notation. Let `%F64` be a notation scope in which `*` and `+` and are interpreted as double-precision (64-bit) floating-point operators and `5.7` is interpreted as a double-precision floating-point constant; let R be the function that injects a floating-point value into the reals. Then we might prove that

$$\frac{1 \leq x \leq 100 \quad -1000 \leq y \leq 1000}{|(x * 5.7 + y) \% F64 - (R(x) \times 5.7 + R(y))| \leq A}$$

where A is an accuracy bound calculated by VCFloat.

When you prove the correctness and accuracy of a numerical program, there is far more to do than bound the round-off error. If we view `(x*5.7+y)%F64` as a *floating-point functional model* of your program, and $R(x) \times 5.7 + R(y)$ as a *real-valued functional model* of the same program, then the main result of interest can be proved by composing these three theorems:

1. The real-valued functional model finds a solution to the mathematical problem of interest, within accuracy bound A_1 .
2. The float-valued functional model approximates the real-valued functional model within accuracy bound A_2 .
3. The program (in C or Python or whatever) correctly implements the float-valued functional model.

VCFloat2 provides

- A *modelling language* for describing float-valued functional models and automatically deriving the corresponding real-valued models;
- A prover for bounding roundoff error, the difference between the two models;
- Tools for connecting float-valued models to C programs. But in VCFloat 2.0 (unlike in 1.0) the float-valued modeling language is quite independent of C and can be used to reason about numerical programs in other languages.

To use VCFloat you follow these steps, we will explain below:

1. Write down your floating-point functional model as Coq functions on floating-point values.
2. Pick identifiers for your variables and apply a tactic to reify your model.
3. Automatically derive a real-valued functional model.
4. Specify bounds for your input variables, in a `boundsmap`.

5. State a roundoff-error theorem, and start the proof using the `prove.rndval` tactic; this leaves two subgoals, “stage 1” and “stage 2”.
6. Prove the stage-1 verification conditions; usually this is as easy as writing `all: interval`.
7. Prove the stage-2 goal; sometimes this is completely automatic, sometimes you have to assist.
8. (optional) Prove that your C program correctly implements the functional model.
9. (optional) Prove that your real-valued model accurately approximates the mathematical quantity of interest.

2 Floating-point functional models

To use VCFLOAT you start with,

Require Import `vcfloat.VCFLOAT`.

This imports VCFLOAT’s functional modelling language and all of its provers.

Functional models are written as expressions in Coq that apply functions (such as *add* and *multiply*) to variables, constants, and subexpressions that belong to floating-point *types*. We will start with the types.

```

type : Type                each floating-point format is described as a type
ftype: type → Type        a floating-point number in format t belongs to Coq type ftype(t)
TYPE : ∀ (precp : positive) (femax: Z), (fprecp < femax) → (1 < fprecp) → type.
Tsingle: type := TYPE 24 128 1 1.
Tdouble: type := TYPE 53 1024 1 1.

```

That is, you specify a floating-point format, a *type*, by the number of mantissa bits (e.g., 24 for single-precision, 53 for double-precision, but any number ≥ 2 is legal) and a maximum exponent value (128 for single-precision, 1024 for double-precision, any number greater than the number of mantissa bits). `TYPE` is a constructor for *type*, and the `1` arguments happen to be proofs that $24 < 128$ and $1 < 24$, and so on.

3 NaNs

In the IEEE-754 floating-point standard, one cannot simply *add* two numbers, one must specify how the NaNs will be propagated. That is, if x and y are double-precision floats, what Not-a-Number (NaN) should float-add return if x or y or both are Not-a-Number? Unfortunately that is left to each computer architecture to decide. VCFLOAT wants to be rigorously faithful to the semantics of the actual computation, so we specify the NaN-propagation behavior of the floating-point model in a typeclass `Nans`.

The good news is that if your computation never produces any NaNs, then it won’t matter which instance of the `Nans` typeclass you use. And VCFLOAT helps you prove that your computation never produces NaNs. Then you can parameterize your float-functional-model over a `NANS` parameter as follows:

Section WITHNANS.
Context {NANS: Nans}.

... your functional model goes here
End WITHNANS.

That is, the NANS variable can be instantiated with *any* architecture-specific which-nans-to-use structure, and then your float-functional-model will consult this NANS structure whenever it produces a not-a-number, which you will prove is never.

4 Notation Scopes

These notation scopes (and their delimiters) come with VCFloat:

Delimit **Scope** float32_scope **with** F32.

Delimit **Scope** float64_scope **with** F64.

Delimiters %F32 and %F64 indicate that constants and operator-symbols should stand for single precision and double-precision (respectively) values and functions.

Definition myformula (h : ftype Tdouble) := (5.0e-1 + cast Tsingle ($h * 1.6$))%F64)%F32.

Here, the constant 1.6 and operator $*$ are interpreted in double precision, and the constant 5.0e-1 (which could just as well have been written as .5) and operator $+$ are interpreted in single precision. The variable h is a double-precision floating-point number.

5 Operators

The following operators are available in each notation scope:

$+$ $-$ $*$ $/$ $<$ $<=$ $>$ $>=$

The minus sign $-$ can be used infix (subtraction) or prefix (negation). The comparison operators can be used in the style $x <= y < z$ as usual in Coq. The following functions can also be used:

BABS (absolute value)

BSQRT (square root)

cast t (cast to ftype(t))

sin, **cos**, **fma**, ... You can also use user-defined functions, as long as you provide proofs (or axioms) about their accuracy; see Section 17.

6 Example

A mass on a spring—a harmonic oscillator—with position x and velocity v can be simulated over time-step $h = \frac{1}{32}$ using the Verlet (“leapfrog”) method with the formula,

Definition h := (1/32)%F32.

Definition $F(x$: ftype Tsingle) : ftype Tsingle := (3.0- x)%F32.

Definition $step$ (x v : ftype Tsingle) := ($x + h*(v+(h/2)*F(x))$)%F32.

Here, the function **step** is the functional model of (part of) a C program:

```

const float h = 1.0/32.0;
float F (float x) { return 3.0f-x; }
float step (float x, float v) { return x+h*(v+h/2.0f)*F(x); }

```

7 Reification

VCFLOAT will *reify* your functional model into the internal syntax tree that it uses.

To represent *in a logic* a function analyzing logical formulas of type τ , one cannot write a function with type $\tau \rightarrow \text{Prop}$; one must operate on *syntactic representations* of floating-point formulas. VCFLOAT's `expr` type is for abstract-syntax trees of formulas (you can do `Print expr` in Coq to see). One can then define in the logic a *reflect* function of type $\text{expr} \rightarrow \tau$. VCFLOAT has `fval` to reflect back to floating-point expressions, and `rval` to reflect into real-valued formulas. (You can do `Check fval` or `Check rval`.)

The opposite process, *reification*, converting from a formula into its abstract-syntax tree, cannot be done *within* the logic, but it can be done by an Ltac program. VCFLOAT provides the tactic `HO_reify_float_expr`. One cannot prove a tactic correct, but you do get a per-instance guarantee for each $f : \tau$ by checking that `reflect(reify(f)) = f`.

The reifier will need a *name* for each of your variables. VCFLOAT's name type is the Coq positive numbers. In our example the variables are x and v , and we will use 1 and 2 for their names:

Definition `_x` : ident := 1%positive. (* Variable name for position *)

Definition `_v` : ident := 2%positive. (* Variable name for velocity *)

Here, the Coq variable `_x` contains not the value, but the *identifier* that we will use for the floating-point variable x . It is not necessary to use consecutive positives, we could have used 5 and 2. Now we can connect `_x` and `_v` to x and v as follows:

Definition `step'` := ltac:(let e' := HO_reify_float_expr constr:([_x; _v]) step in exact e').

This is a tactical definition of a VCFLOAT abstract-syntax tree, `step'`, the reified version of `step`. The tactic is called `HO_reify_float_expr`, and it expects its second argument (in this case, `step`) to be a function from (zero or more) floating-point values to a floating-point value. It learns how many arguments there should be from examining the Coq type of `step`. In this case, since `step` has type `ftype Tsingle → ftype Tsingle → ftype Tsingle`, the tactic knows that `step` should have two arguments, both single-precision floats.

The first argument of `HO_reify_float_expr` should be list of identifiers, to associate with those parameters of the functional model. In this case the list is simply `[_x;_v]`.

8 Boundsmap

In order to do round-off analysis one generally needs *bounds* in the input variables: For example, what are the lowest and highest possible values of x and v in our example? We gather information about each variable (name, floating-point type, low-bound, high-bound) into a **boundsmap**, which maps variable-identifiers to **varinfo** structures.

Record `varinfo` := {var_type: type; var_name: ident; var_lobound: R; var_hibound: R}.

To create the **boundsmap** first make a list of **varinfos**, then use some Ltac boilerplate to compute.

Definition `step_bmap_list` : list varinfo :=

[Build_varinfo Tsingle _x 2 4 ; Build_varinfo Tsingle _v (-2) 2].

Definition `step_bmap` : `boundmap` :=

`ltac:(let z := compute_PTree (boundmap_of_list step_bmap_list) in exact z).`

In the first definition, we make a list of `varinfo` structures. For each parameter of the functional model, we specify its floating-point precision, its identifier, its lowest possible input value, and its highest possible input value. We put these into a list—in our example, `step_bmap_list`. Then the tactical definition (`step_bmap`) is a line of boilerplate that always looks the same (except for the italicized part where you specify this list as shown above).

9 Valmap and reflection

You can *reflect* the abstract-syntax tree (such as `step'`) back into a functional model (such as `step`). To do that, first make a `valmap` that relates your variable identifiers to floating-point values.

Definition `step_vmap_list` ($x\ v : \text{ftype Tsingle}$) := `[(_x, existT ftype _ x);(_v, existT ftype _ v)].`

Definition `step_vmap` ($x\ v : \text{ftype Tsingle}$) : `valmap` :=

`ltac:(let z := ltac:(make_valmap_of_list (step_vmap_list x v)))`.

The auxiliary definition `step_vmap_list` (when applied to x and v) is a list of pairs, identifier \times value, where the “value” is a dependent pair of a type (a floating point format such as `Tsingle` or `Tdouble`) and a value of that type. In this case, both x and v are single-precision, but `valmaps` have the ability to mix precisions.

The second step computes this association list into an efficient data structure.

The function `fval` evaluates the floating-point interpretation of an AST, in an environment that maps the variables. To *reflect* an AST using a `valmap`, apply `fval` as follows:

Definition `reflected_step` ($x\ v : \text{ftype Tsingle}$) := `fval (env_ (step_vmap x v)) step'`.

Lemma `reflect_reify` : $\forall x\ v, \text{reflected_step } x\ v = \text{step } x\ v$.

Proof. `reflexivity. Qed.`

The lemma demonstrates that the round-trip—`reify` then `reflect`—is indeed the identity function.

10 Real-valued functional model

Suppose we take the float-valued functional model (the `step`) function) and interpret every constant and operator in the real numbers:

Definition `step_realmodel'` ($x\ v : \text{ftype Tsingle}$) : `R` := `FT2R x + (1/32)*(FT2R v + ((1/32)/2)*(3- FT2R x))`.

You can make this look prettier using a coercion:

Coercion `FT2R`: `ftype` \rightarrow `R`.

Definition `step_realmodel` ($x\ v : \text{ftype Tsingle}$) : `R` := `x + (1/32)*(v + ((1/32)/2)*(3-x))`.

In fact, you can automatically derive a real-valued functional model using the `rval` function, which reflects into the reals much like `fval` reflects into the floats. Here’s a theorem showing that you get what you’d expect:

Lemma `correspond_floatmodel_realmodel`: $\forall x\ v, \text{rval (env_ (step_vmap } x\ v)) \text{ step'} = \text{step_realmodel } x\ v$.

Proof. `intros. unfold step_realmodel. simpl. repeat f_equal; compute; lra. Qed.`

11 Round-off theorem

The purpose of VCFloat is to prove how accurately the float-valued functional model approximates the real-valued functional model. Here’s an example of such a theorem:

Lemma `prove_roundoff_bound_step`: $\forall vmap, \text{prove_roundoff_bound step_bmap } vmap \text{ step' } (/ 4000000)$.

This says, for any valmap *vmap* containing values for *x* and *v* that are within the bounds specified by `step_bmap`, the difference between the floating-point interpretation of `step'` and the real-number interpretation of `step'` will be less than one four-millionth.

Recall, of course, that “the floating-point interpretation of `step'`” is exactly our float-valued functional model; and “the real-number interpretation of `step'`” is exactly our real-valued functional model.

What if we didn’t know the accuracy $1/400000$ in advance? VCFloat can calculate it; see Section 13.

Here is how we prove the theorem:

Lemma `prove_roundoff_bound_step`: $\forall vmap, \text{prove_roundoff_bound step_bmap } vmap \text{ step' } (/ 4000000)$.

Proof.

`intros.`

`prove_roundoff_bound.`

—

`prove_rndval.` (** see section 12 **)

`all: interval.`

—

`prove_roundoff_bound2.`

`prune_terms (cutoff 30).`

`do_interval.`

Qed.

To prove a `prove_roundoff_bound` theorem, use the `prove_roundoff_bound` tactic. It leaves two subgoals (delimited by “bullets”).

The first subgoal is always proved with `prove_rndval`, which leaves a few verification conditions. In this case, there are three: proving that the additions and subtractions do not overflow. Usually the subgoals left by `prove_rndval` are easy to prove using the Coq Interval package, as shown here by `all: interval`. Section 12 explains and discusses these goals.

The second subgoal is always proved by `prove_roundoff_bound2`, which leaves one subgoal. In this case the subgoal is,

NANS : Nans

$v.v : \mathbb{R}, \quad \text{BOUND} : -2 \leq v.v \leq 2$

$v.x : \mathbb{R}, \quad \text{BOUND0} : 2 \leq v.x \leq 4$

$e0 : \mathbb{R}, \quad E : \text{Rabs } e0 \leq \text{powerRZ } 2 \text{ } (-150)$

$d : \mathbb{R}, \quad E0 : \text{Rabs } d \leq \text{powerRZ } 2 \text{ } (-24)$

$e1 : \mathbb{R}, \quad E1 : \text{Rabs } e1 \leq \text{powerRZ } 2 \text{ } (-150)$

$e2 : \mathbb{R}, \quad E2 : \text{Rabs } e2 \leq \text{powerRZ } 2 \text{ } (-150)$

$d0 : \mathbb{R}, \quad E3 : \text{Rabs } d0 \leq \text{powerRZ } 2 \text{ } (-24)$

$e3 : \mathbb{R}, \quad E4 : \text{Rabs } e3 \leq \text{powerRZ } 2 \text{ } (-150)$

$e4 : \mathbb{R}, \quad E5 : \text{Rabs } e4 \leq \text{powerRZ } 2 \text{ } (-150)$

$d1 : \mathbb{R}, \quad E6 : \text{Rabs } d1 \leq \text{powerRZ } 2 \text{ } (-24)$

----- $(1/1)$

$\text{Rabs } ((v.x + (1/32 * ((v.v + (1/64 * ((3-v.x)*(1+d0)+e1) + e4)) * (1+d) + e3) + e2)) * (1+d1) + e0$
 $- (v.x + 1/32 * (v.v + 1/32 / 2 * (3-v.x))))$

$\leq / 4000000$

That is, the real-valued variables v_x and v_v , which represent the values of x and v , are within the bounds specified in the `boundmap`. The variables $\delta, \delta_0, \delta_1$ that represent relative errors of additions and subtractions, are each less than 2^{-24} in absolute value. The variables $\epsilon, \epsilon_0, \epsilon_1, \epsilon_2, \epsilon_3, \epsilon_4$ that represent absolute errors of additions and multiplications are each less than 2^{-150} in absolute value. Finally, assuming all of that, one must prove that the difference between the computation *with* all the deltas and epsilons and the computation *without* the deltas and epsilons, is less than the accuracy bound.

To prove this, one can use the Coq Interval package. But in many cases one must do some work to prepare the goal for solution by Interval. Later chapters will explain. In this example, we use the `prune_terms` tactic. Then the goal solves by `do_interval`.

12 Verification conditions

The first subgoal left by the `prove_roundoff_bound` tactic takes the form,

`prove_rndval bmap vmap step'`

which encapsulates all the *verification conditions* for evaluating the floating-point expressions. These VCs include that certain subexpressions don't overflow into infinities or NaNs, that certain subexpressions don't underflow into subnormal numbers, or even that certain subexpressions *always* underflow.

The `prove_rndval` tactic expands the `prove_rndval` goal into one subgoal per VC. A typical subgoal might be equivalent to this:

```
2 goals
NANS : Nans
v_v, v_x : ftype Tsingle
BOUND : -2 <= FT2R v_v <= 2
BOUND0 : 2 <= FT2R v_x <= 4
e1, d1, e2, e3, d2, e4 : R
H : Rabs e1 <= powerRZ 2 (-150)
H0 : Rabs d1 <= powerRZ 2 (-24)
H1 : Rabs e2 <= powerRZ 2 (-150)
H2 : Rabs e3 <= powerRZ 2 (-150)
H3 : Rabs d2 <= powerRZ 2 (-24)
H4 : Rabs e4 <= powerRZ 2 (-150)
-----
0 < powerRZ 2 128 - Rabs (((FT2R v_x + (powerRZ 2 (-5) * FT2R v_v + e4)) * (1 + d2)
+ e3 + (powerRZ 2 (-11) * - FT2R v_x + e2)) * (1 + d1) + e1)
```

This is basically a proof that the floating point evaluation of our formula on x and v does not overflow. That is, $0 < 2^{128} - |E|$, where E is some function of x and v . E has the same order of magnitude as x and v —both of which are bounded by 4. So it should be *extremely easy* to prove this inequality, and indeed, just calling the `interval` tactic does the job. It's often the case that *all* the subgoals of the `prove_rndval` tactic are immediately provable this way.

13 Letting VCFloat calculate the accuracy bound

You can ask VCFloat not only to *prove* an accuracy bound, but to *calculate and prove* the bound if you don't know it in advance.


```

Derive step_b
SuchThat (∀ vmap, prove_roundoff_bound step_bmap vmap step' step_b)
As prove_step_bound.
Proof. ... Qed.

```

Here we are using Coq's `Derive` command (look it up in the Coq manual). This gives the following initial proof goal:

```

NANS : Nans
step_b := ?Goal : R
-----(1/1)
∀ vmap : valmap,
prove_roundoff_bound step_bmap vmap step' step_b

```

To prove this, simply `subst step_b` to put the unification variable `?Goal` below the line, and proceed with exactly the same proof as shown at `prove_roundoff_bound_step`. Then `do_interval` will instantiate the bound.

After the **Qed**, there will be two new things in the environment:

```

step_b : R = 4644337115725828 / 18889465931478580854784 +
          5910977010728984 * / 9671406556917033397649408
prove_step_bound: ∀ vmap: valmap, prove_roundoff_bound step_bmap vmap step' step_b

```

Displaying bounds in floating point. The bounds derived by the Interval package (via our `do_interval` tactic) will typically be constructed from rational arithmetic, as shown in the `step_b` example here. To see this as a floating-point number, you can use the `ShowBound` tactic.

```

Check ltac:(ShowBound step_b).
(* 2.464803403965517e-7%F64 : Bits.binary64 *)

```

14 field_simplify

The subgoal remaining after `prove_roundoff_bound2` has the form, $|E' - E| \leq A$, where E is the real-valued functional model, E' is the real-valued model with deltas and epsilons inserted to represent round-off errors, and A is the accuracy bound. Goals of this kind can be given to the Coq Interval package to solve. Coq.Interval works by computing in interval arithmetic, where every number is represented by a pair of floating-point numbers representing a lower bound and upper bound. For example, to subtract $a - b$ represented as $(a_{lo}, a_{hi}) - (b_{lo}, b_{hi})$, the result is $(a_{lo} - b_{hi}, a_{hi} - b_{lo})$. For soundness Interval even takes care to set the floating-point rounding modes to round down when computing the lo-bound, and round-up when computing the high bound.

When variables such as `v_x` and `v_v` can take on a wide range of possible values, Interval uses repeated *bisection* to measure many subranges of `v_x` and `v_v`, taking the maximum error.

But there's a problem. Consider the interval calculation of $(a + \delta) - a$, which comes out to

$$(a_{lo} + \delta_{lo} - a_{hi}, a_{hi} + \delta_{lo} - a_{lo}).$$

If the value of a is only approximately known, so $a_{hi} - a_{lo}$ is large, then the interval approximation of $(a + \delta) - a$ is similarly large—which is bad. But we know that whatever the *true* value of a is, when subtracted from itself it will yield zero. That is, a much better approximation can be obtained by *symbolically* subtracting $a + \delta - a = \delta$ before calling the interval package.

If you examine the proof goal at the end of section 11, you'll see that it contains (more or less) $v_x(1 + \delta) - v_x$ and $v_v(1 + \delta) - v_v$. So we can expect the Interval tactic to perform badly on this expression. The solution is to symbolically simplify the expression, and Coq's `field_simplify` tactic can do that:

match goal with \vdash `Rabs ?a <= _` \Rightarrow `field_simplify a` **end**.

This changes the below-the-line portion of the proof goal to,

```
Rabs((-v_x*d0*d*d1-v_x*d0*d-v_x*d0*d1-v_x*d0-v_x*d*d1-v_x*d+2047*v_x*d1+64*v_v*d*d1+
64*v_v*d+64*v_v*d1+3*d0*d*d1+3*d0*d+3*d0*d1+3*d0+e1*d*d1+e1*d+e1*d1+e1+64*e4*d*d1+
64*e4*d+64*e4*d1+64*e4+3*d*d1+3*d+64*e3*d1+64*e3+2048*e2*d1+2048*e2+3*d1+2048*e0)/2048)
<= / 4000000
```

Applying the interval tactic solves this goal immediately.

But `field_simplify` is not the best tool for this job; we use it here only to illustrate the principle of (automatically) expanding the formula into a multinomial and symbolically canceling terms. There are two problems: (1) the multinomial can have an exponential number of terms, most of which are negligible; and (2) floating-point functional models do not always expand into nice multinomials.

15 prune_terms

The proof goal at the end of the last section had many terms similar to $3*d0*d*d1$ where two or more deltas or epsilons are multiplied together. Since the deltas are bounded by 2^{-24} (in single precision) or 2^{-53} (in double precision), and the epsilons are much smaller than that, their product is probably negligible.

The `prune_terms` tactic expands a formula into a multinomial (like `field_simplify`) but also deletes (and bounds) negligible terms—you specify a “cutoff” for what you consider negligible. For example, at the section-11 proof goal, one can write,

`prune_terms (cutoff 30).`

which expands into a multinomial, cancels terms symbolically, and deletes all terms that can be bounded by 2^{-30} . The result is,

```
Rabs (1 * d1 * v_x + 1/32 * d * v_v + 1/32 * d1 * v_v)
<= / 4000000 - 5910977010729000 / 9671406556917033397649408
```

In this goal, the number $5910977010729000 / 9671406556917033397649408$ is the sum of the bounds of the negligible terms. The goal solves easily by the interval tactic.

For comparison, using `(cutoff 50)` gives the goal,

```
Rabs (-1/2048 * d0 * v_x + -1/2048 * d * v_x + 2047/2048 * d1 * v_x +
1/32 * d * v_v + 1/32 * d1 * v_v + 3/2048 * d0 + 3/2048 * d + 3/2048 * d1)
<= / 4000000 - 5242471455916076 / 20282409603651670423947251286016
```

in which not as many terms have been neglected. But either way, `interval` solves the goal.

16 error_rewrites

Suppose an expression for the absolute forward error does not expand into a nice multinomial that is tractable for the `prune_terms` tactic; suppose it is such a large expression that applying the

`field_simplify+interval` tactic causes Coq to crash. An example of such a problem is the `carbonGas` benchmark from the `FPBench` benchmark suite:

`carbonGas(v) := P + A * (N/v) * (N/v) * (v - N * B) - K * N * T`

where `A, B, K, N, P, T` are constants. Because `v` appears in the denominator of terms in `carbonGas`, `VCFLOAT`'s `prune_terms` tactic won't simplify the expression enough for `interval` to produce a decent bound. A tactic that can be used in cases like this is `error_rewrites`, which will recursively decompose a main proof goal for absolute forward error into subgoals of smaller subexpressions on related terms using the following equalities.

$$\begin{aligned}(\tilde{u} - \tilde{v}) - (u - v) &= (\tilde{u} - u) - (\tilde{v} - v) \\(\tilde{u} + \tilde{v}) - (u + v) &= (\tilde{u} - u) + (\tilde{v} - v) \\(\tilde{u} * \tilde{v}) - (u * v) &= (\tilde{u} - u) * v + (\tilde{v} - v) * u + (\tilde{u} - u) * (\tilde{v} - v) \\ \frac{u'}{v'} - \frac{u}{v} &= (u' - u) - (v' - v) * \frac{1}{v} * u\end{aligned}$$

We write \tilde{e} to denote an expression with deltas (for relative error) and epsilons (for absolute error); so $\tilde{u} - u$ is just the absolute error in computing the formula u in floating-point (rather than in the real numbers).

Consider the expression for the absolute forward error of `carbonGas`, which looks like

$$\text{Rabs}((\tilde{u} * (1 + \delta_2) + \epsilon_7 - \tilde{v}) * (1 + \delta_6) + \epsilon_0 - (u - v)).$$

Applying `error_rewrites` produces – supposing that \tilde{u} , \tilde{v} , u , and v are opaque – three subgoals:

subgoal 1: $\text{Rabs}(\tilde{u} * (1 + \delta_2) * (1 + \delta_6) + \epsilon_7 - u) \leq ?e3$

subgoal 2: $\text{Rabs}(\tilde{v} * (1 + \delta_6) - v) \leq ?e2$

subgoal 3: $\text{Rabs}(\epsilon_0) \leq ?e1$

where $?e1, ?e2, ?e3$ are unification variables to be determined in subproofs, and the total error is now bounded by $e3 + e2 + e1$.

The subexpressions in the above subgoals are smaller and contain related terms. For these reasons, using the `field_simplify+interval` tactic is potentially more tractable. However, in some cases, the recursion stops before it fully decomposes a subexpression into a form that `interval` can provide a decent bound on. In this case, as long as all rational expressions have been decently decomposed by the division case in `error_rewrites`, the `prune_terms` tactic can be used successfully.

There are some cases when `error_rewrites` causes Coq to crash: on expressions with a large number of operations (approximately 40). This occurs because `error_rewrites` simply produces too many subgoals. As an example, consider that the `carbonGas` benchmark has 11 floating-point operations and that `error_rewrites` produces 125 subgoals (and this is even after `error_rewrites` has automatically discharged a few!).

17 User-defined functions

Any function that takes 0 or more float arguments (not necessarily all of the same type) and produces a float result can be used in VCFloat. You must provide a `floatfunc_package` to describe its characteristics:

```
Record floatfunc (args: list type) (result: type)
  (precond: klist bounds args)
  (realfunc: function_type (map RR args) R) :=
{ff_func: function_type (map ftype' args) (ftype' result);
 ff_rel: N;
 ff_abs: N;
 ff_acc: acc_prop args result ff_rel ff_abs precondition realfunc ff_func
}.
```

```
Record floatfunc_package (ty: type) :=
{ff_args: list type;
 ff_precond: klist bounds ff_args;
 ff_realfunc: function_type (map RR ff_args) R;
 ff_ff: floatfunc ff_args ty ff_precond ff_realfunc}.
```

Suppose p is floatfunc package, that is, $p: \text{floatfunc_package}(\text{ty})$. Then `ff_args p` are a list of argument types, and `ty` is the result type. For example, suppose `ff_args p` = `[Tdouble, Tsingle]` and `ty` = `Tsingle`, then the function $f = \text{ff_func}(\text{ff_ff } p)$ has type $f: \text{ftype Tdouble} \rightarrow \text{ftype Tsingle} \rightarrow \text{ftype Tsingle}$, because `function_type (map ftype' (ff_args p)) (ftype' ty)` is convertible to `ftype Tdouble \rightarrow ftype Tsingle \rightarrow ftype Tsingle`.

The specification of this function is as follows: There is a real-valued function `ff_realfunc(p): $\mathbb{R} \rightarrow \mathbb{R} \rightarrow \mathbb{R}$` ; and `ff_acc` is the theorem that f approximates this function within relative error `ff_rel p` and absolute error `ff_abs p` , provided that the arguments to f are within the bounds specified by `ff_precond`.

The development `Test/TestFunc.v` in the VCFloat repository demonstrates how to use floatfunc packages.

18 Abstract *versus* transparent valmap

When you have proved a `prove_roundoff_bound` theorem as described in the previous sections, then you may want to use that in the proof of other theorems about your floating-point program. Since `prove_roundoff_bound` quantifies over all valmaps, you can apply it to a particular valmap.

For example, `step_vmap` defined in Section 9 takes floating-point values x and v , and produces a valmap. That is, `step_vmap 3.1 0.7` is the valmap with position 3.1 and velocity 0.7. You could apply the `prove_roundoff_bound_step` theorem to this valmap to prove that, when the program is run with arguments 3.1 and 0.7, then its roundoff error will be less than $1/4000000$.

In that theorem, the *vmap* is a quantified variable—it is abstract. But you can also prove a `prove_roundoff_bound` theorem with a less-abstract vmap.

For example, a theorem like this one (and see `Test/summation.v` for another example):

```
Lemma prove_rndoff' :
   $\forall (x \ v : \text{ftype Tsingle}),$ 
  let accuracy := some function of  $x$  and  $v$ 
  my_extra_constraint  $x \ v \rightarrow$ 
  prove_roundoff_bound step_bmap (step_vmap  $x \ v$ ) step' accuracy.
```

This example illustrates another thing as well: Suppose you have extra constraints on your variables, besides just the boundsmap. You can state these constraints as an extra hypothesis in the theorem, and the proof tactics (such as `interval`) can make use of them. You can also prove a theorem in which the accuracy bound is dependent on the variables in your valmap—for example, a *relative error* bound.

19 Annotations

Floating-point error analysis can be slightly more precise in certain cases:

Denorm: When the result of a calculation is known to be a *denormal* (also called *subnormal*) number—a tiny number within $2^{e_{\min}}$ of zero—then it has only an additive error. That is, $(a + b) + \epsilon$ instead of $(a + b)(1 + \delta) + \epsilon$.

Norm: When the result of a calculation is known to be a *normal* number—that is, bounded away from zero by at least $2^{e_{\min}}$ —then it has only a relative error. That is, $(a + b)(1 + \delta)$ instead of $(a + b)(1 + \delta) + \epsilon$.

Sterbenz: When a, b satisfy $\frac{1}{2} < \frac{a}{b} < 2$, then the floating point subtraction $a - b$ is exact, no relative error δ , no absolute error ϵ .

You can annotate these cases in your functional model using these functions:

Definition `Norm {A}(x: A) := x.`

Definition `Denorm {A}(x: A) := x.`

Definition `Sterbenz {A}(x: A) := x.`

As you can see, these are just identity functions, so semantically they do nothing. But they guide VCFLOAT’s reifier to mark its internal abstract-syntax tree. This will cause additional proof obligations (subgoals) at stage 1, to *prove* that such-and-such a subexpression is normal, or denormal, or Sterbenz; but will cause fewer deltas and epsilons to be generated at stage 2.

In our running example we could write,

Definition `h := (1/32)%F32.`

Definition `F(x: ftype Tsingle) : ftype Tsingle := Sterbenz(3.0-x)%F32.`

Definition `step (x v: ftype Tsingle) := Norm(x + h*(v+(h/2)*F(x)))%F32.`

20 Verified Software Toolchain

You can use the Verified Software Toolchain (VST) to prove that a C program correctly implements a floating-point functional model.

Along with importing `VST.floyd.proofauto` and the other standard boilerplate that introduces a VST proof, you will want:

From `vcfloat` **Require Import** `FPCompCert Float_notations.`

Require Import `float_model. (* your functional model *)`

It is not necessary to import all of `vcfloat.VCFLOAT`; the imports shown are enough to connect CompCert’s definitions for floating point to VCFLOAT’s definitions. CompCert and VCFLOAT use different names for the same underlying Flocq floating-point types:

```

Eval compute in compcert.lib.Floats.float32. (* = Binary.binary_float 24 128 *)
Eval compute in ftype Tsingle.               (* = Binary.binary_float 24 128 *)
Eval compute in compcert.lib.Floats.float.    (* = Binary.binary_float 53 1024 *)
Eval compute in ftype Tdouble.                (* = Binary.binary_float 53 1024 *)

```

In your VST assertions (funspecs, loop invariants, etc.), use the VCFloat names for those types. For example, here we write `ftype Tsingle` instead of `float32`:

```

Definition force_spec :=
  DECLARE _force
  WITH q : ftype Tsingle
  PRE [ tfloat ] PROP() PARAMS(Vsingle q) SEP()
  POST [ tfloat ] PROP() RETURN (Vsingle (F q)) SEP().

```

There are two useful tactics to convert CompCert-style float notations to VCFloat-style notations. Immediately after `start.function` you can write,

```

start.function.
subst MORE_COMMANDS; unfold abbreviate; canonicalize_float_constants.

```

The tactic `canonicalize_float_constants` converts all of the floating-point literals in the AST of your function-body into a VCFloat style, which makes them easier to reason about (and to relate to your functional model).

The following tactic is useful *after* going forward through a sequence of C statements that perform floating-point operations:

```

autorewrite with float_elim in *.

```

It converts `Float32.add x y` to `(x+y)%F32`, and similarly for other operators.

That’s it! Other than these conversions, you use VST in a completely standard way.

Exactly matching the functional model

When using VST or any other tool to prove that a program correctly implements a functional model, take care to match it exactly. For example, in floating point the associative law $(a+b)+c = a+(b+c)$ does not hold.

The commutative law $a+b = b+a$ may hold *only if neither a nor b is a NaN*, depending on how your target machine propagates NaNs. Therefore, if your C program computes $a+b$ while the functional model computes $b+a$, you will be able to prove it correct only if you propagate the invariant that a is finite and b is finite. It’s certainly possible to propagate such invariants, but it’s simpler if you don’t have to.

21 Examples

VCFloat comes with several worked examples, in the `Test` and `FPBench` directories:

Test/TestRefman.v the running example from this reference manual

Test/TestPaper.v the (similar) running example from “VCFloat2: Floating-point error analysis in Coq”

Test/TestFunc.v An example of user-supplied functions.

FPBench/*.v examples from the FPBench benchmark suite (fpbench.org)

22 Nonstandard floating-point-like types

Users may implement in software floating-point types that do not correspond exactly to any IEEE 754 format but that can be shown to respect round-off error bounds. VCFloat2 supports such user-defined types.

Consider the example of double-double [?]. A double-double number can be used to represent a floating-point number with at least 106 bits of mantissa using two double-precision floats whose 53-bit mantissas do not overlap (because their exponents differ by at least 53). One can add or multiply numbers in this data type using just a few ordinary double-precision operations, especially on machines that support fused multiply-add (fma).

The rounding behavior of double-double can be shown to be follow model with relative error $|\delta| \leq 2^{-106}$ and absolute error $|\epsilon| \leq 2^{-1022}$. If the user can prove such a property of an abstract number type (of which double-double is just one example), then they can build a `nonstdtype` record in VCFloat2:

Record `nonstdtype`

```
(fprec: Z) (femax: Z) (prec_range: 1 < fprec < femax) :=
NONSTD
{ nonstd_rep: Type;
  nonstd_to_F: nonstd_rep → option (float radix2);
  ... (* some fields omitted *) ...
  nonstd_bounds: ∀ x: nonstd_rep,
    ( - (bpow radix2 femax - bpow radix2 (femax - fprec)) <=
      floatopt_to_real (nonstd_to_F x) <=
        bpow radix2 femax - bpow radix2 (femax - fprec) ) %R
}.
```

Here, `nonstd_bounds` is the user-supplied proof of such a rounding theorem, `nonstd_to_F` is the user-supplied definition of a function that maps nonstandard types to a floating-point representation, and `floatopt_to_real` is a function provided by VCFloat2 that maps a floating-point representation to a real number.

After building a `nonstdtype` record, the `GTYPE` constructor can be used to build a new VCFloat2 type (see Chapter 23). Expressions in which some functions return values of this type and other functions take values of this type can now be reified, and VCFloat2 will automatically calculate and prove round-off error bounds.

In order to write such expressions, users of VCFloat2 must first define operations on the newly defined type. This is done by constructing a `floatfunc` record (whose type was presented in Section ??).

Standard vs. nonstandard types. One might wonder what is special about a “standard” type, that it cannot be just an instance of the general notion of nonstandard type. A standard type must support all the standard binary and unary operations (`Zconst`, `BINOP`, `UNOP`, `PLUS`, `MINUS`, ... see Chapter 26). Nonstandard types support only whatever functions someone has supplied. The notion of an arbitrary “cast” from any IEEE precision to any other cannot be generically supported, and (for example) some nonstandard types do not support Sterbenz subtraction, or constant literals.

23 Functional Modeling Languages

VCFloat comes with (two variants of) a *functional modeling language* for describing floating-point computations as functional programs. Chapters 2–5 gave an informal introduction. The modeling language(s) are useful not only for calling VCFloat’s automated roundoff analysis but for other reasoning.

The *General* modeling language is VCFloat’s default: it permits IEEE-754 floating-point types at any size of exponent and mantissa, as well as *nonstandard* types such as flush-to-zero or double-double. The *Standard* modeling language permits only the IEEE-754 floating-point types at any size of exponent and mantissa; for some applications it is more convenient to use.

<i>Language</i>	General	Standard
<i>Purpose 1</i>	Ordinary use of VCFloat’s automated roundoff calculations	(can be used, with conversion, for automated roundoff calculations)
<i>Purpose 2</i>	Modeling computations that may involve nonstandard floats	Modeling computations using only IEEE floats
<i>Advantage</i>	Generality	Easier conversion between <code>ftype(t)</code> and <code>binary_float(fprec t)(femax t)</code>
<i>Semantics</i>	Record type: <code>Type :=</code> <code>GTYPE {fprec: Z; femax: Z; ...;</code> <code>nonstd: option (nonstdtype ...)}.</code> Definition <code>TYPE fprec femax ... :=</code> <code>GTYPE fprec femax ... None.</code>	Record type: <code>Type :=</code> <code>TYPE {fprec: Z; femax:Z, ...}</code>
<i>Import</i>	Require Import vcfloat.VCFloat.	From vcfloat Require vcfloat.VCFloat. Require Import vcfloat.FPStdLib.
<i>Alt. Import</i>	Require Import vcFloat.FPLib. <i>import this way when you don’t need automatic roundoff calculation</i>	From vcfloat Require RAux FPStdLib.
<i>Alt. Import</i>		Require Import vcfloat.FPStdLib. <i>when you don’t need much reasoning about how floats represent reals</i>

Aside from automatic roundoff calculations, both modeling languages come with several useful libraries for reasoning about floating-point:

Chapter 26: Floating-point expressions that are generic over their precision, in an easier way than basic Flocq provides.

Chapter 4: Notation systems for floating-point literals and expressions.

Chapter 27: Equivalence relations on floating-point numbers, with Coq *Morphisms* for automatic rewriting.

Chapter 28: Automatic generation of real-valued functional models from float-valued models.

24 General Modeling Language

The default functional modeling language in VCFLOAT is based on the notion of a floating-point format, called a type:

Record type: `Type := GTYPE`

```
{ fprec: positive;
  femax: Z;
  fprec := Z.pos fprec;
  fprec_lt_femax_bool: ZLT fprec femax;
  fprec_not_one_bool: Bool.is_true (negb (Pos.eqb fprec xH));
  nonstd: option (nonstdtype fprec femax fprec_lt_femax_bool fprec_not_one_bool)
}.
```

The number `fprec` is the number of bits in the mantissa, `femax` is the largest possible binary exponent. Propositions `fprec_lt_femax_bool` and `fprec_not_one_bool` guarantee that $1 < \text{fprec} < \text{femax}$. Finally, `nonstd` is `None` for an ordinary IEEE type (even of a nonstandard mantissa size or exponent size), and `Some` for truly weird formats such as flush-to-zero¹ or double-double.

The constructor `GTYPE` builds a type, and the derived function `TYPE` builds a standard type:

Definition `TYPE fprec femax fprec_lt_femax fprec_not_one :=`
`GTYPE fprec femax fprec_lt_femax fprec_not_one None.`

Definition `Tsingle := TYPE 24 128 I I. (* IEEE 32-bit *)`

Definition `Tdouble := TYPE 53 1024 I I. (* IEEE 64-bit *)`

Floating point numbers in format t belong to (in Coq's type system) to `ftype(t)`. One can convert between `ftype(t)` and the Flocq representation of IEEE types by using rewriting:

Definition `ftype: type → Type := ...`

Class `is_standard (t: type) :=`

`is_standard: match nonstd t with None ⇒ True | _ ⇒ False end.`

Definition `ftype_of_float {t: type} {STD: is_standard t}:`
`binary_float (fprec t) (femax t) → ftype t.`

Definition `float_of_ftype {t: type} {STD: is_standard t}:`
`ftype t → binary_float (fprec t) (femax t).`

Lemma `float_of_ftype_of_float {t: type} (STD1 STD2: is_standard t):`
 $\forall x, \text{@float_of_ftype } t \text{ STD1 } (\text{@ftype_of_float } t \text{ STD2 } x) = x.$

Lemma `ftype_of_float_of_ftype {t: type} (STD1 STD2: is_standard t):`
 $\forall x, \text{@ftype_of_float } t \text{ STD1 } (\text{@float_of_ftype } t \text{ STD2 } x) = x.$

In contrast, in the *Standard* modeling language (described in the next chapter), `ftype t` is convertible with `binary_float (fprec t) (femax t)` by *unification*, without needing rewriting; this convenience is the main reason for having the *Standard* language.

¹Some GPUs lack subnormal numbers, so instead of gradual underflow when one reaches the minimum (most negative) exponent, they round directly to zero.

25 Standard Modeling Language

The nondefault “Standard” modeling language based on the notion of a floating-point format, called a *type*, but *without a nonstd option*.

The default functional modeling language in VCFLOAT is based on the notion of a floating-point format, called a *type*:

```
Record type: Type := TYPE
{ fprec: positive;
  femax: Z;
  fprec := Z.pos fprec;
  fprec_lt_femax_bool: ZLT fprec femax;
  fprec_not_one_bool: Bool.is_true (negb (Pos.eqb fprec xH))
}.
```

Definition ftype ty := binary_float (fprec ty) (femax ty).

Definition Tsingle := TYPE 24 128 | I. (* IEEE 32-bit *)

Definition Tdouble := TYPE 53 1024 | I. (* IEEE 64-bit *)

The number *fprec* is the number of bits in the matissa, *femax* is the largest possible binary exponent. Clearly, *ftype(t)* is definitionally equal (hence unifiable) with Flocq’s *binary_float (fprec t) (femax t)*.

One can **Require** both modeling languages at once, but it is recommended not to **Import** both at once, since they disagree about the meaning of *type*, *ftype*, *Tsingle*, *et cetera*.

26 Generic expressions

Consider the harmonic oscillator example shown in Chapter 6. That was given in single-precision floating point, but one could model the algorithm at arbitrary precision, and do proofs about it. We would write,

Require Import vcfloat.FPLib.

Section WITHNANS.

Context {NANS: Nans}.

Definition h {t: type} '{STD: is_standard t} := BDIV (Zconst t 1) (Zconst t 32).

Definition F {t: type} (x: ftype t) '{STD: is_standard t} : ftype t := BMINUS (Zconst t 3) x.

Definition step {t: type} '{STD: is_standard t} (x v: ftype t) :=

BPLUS x (BMULT h (BPLUS v (BMULT (BDIV h (Zconst t 2)) (F x)))).

End WITHNANS.

Instead of importing FPLib (which gives access only to the “General” modeling language and morphisms) one could import all of vcfloat.VCFLOAT (which gives access in addition to the roundoff-error automation). However, roundoff-error automation works only at specific precisions; so after defining *step* in this way, one could specialize it to *@step Tsingle _* or *@step Tdouble _* and call the automation. However, there are many uses of VCFLOAT’s supporting libraries even when its roundoff-error automation is not used.²

²Here are two examples. LAPROOF: a library of formal accuracy and correctness proofs for sparse linear algebra programs, by Ariel E. Kellison, Andrew W. Appel, Mohit Tekriwal, and David Bindel, 30th IEEE International Symposium on Computer Arithmetic, September 2023 (github.com/VeriNum/iterative_methods). Verified correctness, accuracy, and convergence of a stationary iterative linear solver: Jacobi method, by Mohit Tekriwal, Andrew W. Appel, Ariel E. Kellison, David Bindel, and Jean-Baptiste Jeannin, 16th Conference on Intelligent Computer Mathematics, September 2023 (github.com/VeriNum/LAPROOF).

If using FPStdLib instead of FPLib, to use the “Standard” modeling language, then omit the STD arguments from the functions.

	General	Standard
3	Zconst (t: type) '{STD: is_standard t} (i: Z) : ftype t	Zconst (t: type) (i: Z) : ftype t
	BINOP (op) (nan) {t: type} '{STD: is_standard t}: ftype t → ftype t → ftype t	BINOP (op) (nan) {t: type}: ftype t → ftype t → ftype t
+	BPLUS := BINOP Bplus plus_nan	BPLUS := BINOP Bplus plus_nan
−	BMINUS := BINOP Bminus minus_nan	BMINUS := BINOP Bminus minus_nan
	UNOP (op) (nan) {t: type} '{STD: is_standard t}: ftype t → ftype t	UNOP (op) (nan) {t: type}: ftype t → ftype t
·	BABS {t: type} '{STD: is_standard t}	BABS {t: type}
−	BOPP {t: type} '{STD: is_standard t}	BOPP {t: type}
√	BSQRT := UNOP Bsqr sqrt_nan	BSQRT := UNOP Bsqr sqrt_nan
− · − + −	BFMA (<i>fused multiply-add</i>)	BFMA
	The nan parameter gives instructions on forming a NaN payload. BOPP, BABS cannot produce NaNs from non-NaNs, so do not need this parameter.	

27 Equivalence relations and rewriting morphisms

When is the expression $a \cdot 0 + b$ equal to b ? This is important, for example, in modeling sparse matrix operations, where a dense representation computes the $a \cdot 0 + b$ where the sparse representation has simply b .

Well, if $b = 3.5$ and $a = \infty$ or $a = \text{NaN}$, then $a \cdot 0 + b = \text{NaN}$. So we might ask, if a and b are known to be finite, is $a \cdot 0 + b$ equal to b ? And unfortunately, if a is *negative zero* and b is *positive zero*, then the fused-multiply-add $a \cdot 0 + b = -0$ while $b = +0$.

This may be of little consequence when reasoning in the reals, where $-0 = +0$. But often, in proving that a low-level implementation (in C or some other language) correctly implements a floating-point functional model, we want to avoid reasoning about the real-valued *meaning* of the terms, and just treat floating-point operations as uninterpreted functions. But even in that reasoning we want to consider $-0 \cong +0$.

For this purpose, `vcfloat.FPLib` and `vcfloat.FPStdLib` (choose one) provide these relations:

`feq x y` means that either x and y are both nonfinite (infinities or NaNs), or they are both zeros (regardless of sign), or they are identical.

`strict_feq x y` means x and y are both finite, and either they are both zeros (regardless of sign) or they are identical. This is a *partial equivalence relation* (or “setoid”).

These relations are theorized in Coq’s Setoid Morphism system, so you can use them in rewriting. (There is a third relation, *mathsf{float}_equiv x y*, meaning that either x and y are both NaNs or they are identical, which VCFLOAT uses internally but which is less useful for general purposes and is not theorized as a setoid.)

For example, if `feq x x' ∧ feq y y' ∧ feq z z'` then `feq (x · y + z) (x' · y' + z')`, and one can prove this easily by rewriting.

28 Conversions between models

Given a floating-point functional model, one can convert it into the corresponding real-valued functional model by reifying and then using the `rval` function, as described in Chapter 10.

Given a model in the *Standard* modeling language, one can automatically convert it into the *Generic* modeling language using the `type_coretype` relation. But this is experimental and rather ill-documented.

29 Bibliography

VCFloat 1.0 was built in 2015 and described in,

A unified Coq framework for verifying C programs with floating-point computations, by Tahina Ramananandro, Paul Mountcastle, Benoît Meister, and Richard Lethin, in *Proceedings of the 5th ACM SIGPLAN Conference on Certified Programs and Proofs (CPP'16)*, pages 15–26, 2016 (<https://doi.org/10.1145/2854065.2854066>).

VCFloat 2.2 was built 2021-2023 and described in,

VCFloat2: Floating-point error analysis in Coq, by Andrew W. Appel and Ariel E. Kellison, in *CPP'24: Proceedings of the 13th ACM SIGPLAN International Conference on Certified Programs and Proofs*, January 2024 (distributed as `doc/vcfloat2.pdf` in the `vcfloat` repo).

VCFloat 2.0 was applied and demonstrated in,

Verified numerical methods for ordinary differential equations, by Ariel E. Kellison and Andrew W. Appel, in *15th International Workshop on Numerical Software Verification (NSV'22)*, August 2022.